# DataDough 🦥

DataDough is a comprehensive analysis of publicly available data sourced from kaggle, conducted by Shreya Sudan. It encompasses a series of rigorous analytical procedures, including data cleansing, exploratory data analysis, data visualizations, hypothesis testing, and predictive analysis

In [1]:
import pandas as pd
import numpy as np
import plotly.express as px

### Introduction 99



Welcome to DataDough, a data science project focused on exploring the fascinating realm of data and uncovering insights related to salaries of data scientists. In today's data-driven world, understanding the factors that influence data scientists' salaries is crucial for both professionals and organizations alike. This project aims to shed light on the various factors that contribute to the salaries of data scientists and provide valuable insights into this ever-evolving field.

The dataset salary contains data on not only the salaries of data scientists, but also their experience, salary in USD, residence, company's location, and company's size.

This project not only seeks to answer questions about salary trends but also aims to explore the relationships between different variables and their impact on compensation. By examining the interplay between factors like work experience, employment type, and company size, we aim to uncover valuable insights that can aid professionals in negotiating salaries and assist organizations in formulating

To accomplish these objectives, we employ a range of analytical techniques and tools, including data preprocessing, exploratory data analysis, feature engineering, and machine learning algorithms. The project utilizes popular libraries like pandas, scikit-learn, and plotly to perform data manipulation, modeling, and visualization.

Here are the components of our project:

- Part I : Inferential Analysis
  - Data Cleaning
  - Exploratory Data Analysis (EDA)
  - Hypothesis Testing
- Part II : Predictive Analysis
  - Linear Regression Model
    - o Framing the Problem
    - o Baseline Model
    - Final Model
  - Multiclass Classification
    - o Random Forest Classifier
    - o Decision Trees

By the end of DataDough, we aim to provide a comprehensive understanding of the salary landscape for data scientists, empowering both individuals and organizations with actionable insights. Whether you are a data scientist seeking to benchmark your salary or an organization looking to attract and retain top talent, this project strives to offer valuable information and assist in data-driven decision-making

In [2]: salary = pd.read\_csv('ds\_salaries.csv')

# print(salary.head().to\_markdown(index=False))

t[2]:	work_y	ar experience_leve	el employment_type	job_title	salary	salary_currency	salary_in_usd	employee_residence	remote_ratio	company_location	company_size
C	20	23 S	E FT	Principal Data Scientist	80000	EUR	85847	ES	100	ES	L
1	20	23 N	II CT	ML Engineer	30000	USD	30000	US	100	US	S
2	20	23 N	II CT	ML Engineer	25500	USD	25500	US	100	US	s
3	20	23 S	E FT	Data Scientist	175000	USD	175000	CA	100	CA	М
4	20	23 S	E FT	Data Scientist	120000	USD	120000	CA	100	CA	М

In [3]: salary.describe()

Out[3]:

# print(salary.describe().to\_markdown())

		work_year	salary	salary_in_usd	remote_ratio
cou	unt	3755.000000	3.755000e+03	3755.000000	3755.000000
me	an	2022.373635	1.906956e+05	137570.389880	46.271638
5	std	0.691448	6.716765e+05	63055.625278	48.589050
n	nin	2020.000000	6.000000e+03	5132.000000	0.000000
25	5%	2022.000000	1.000000e+05	95000.000000	0.000000
50	)%	2022.000000	1.380000e+05	135000.000000	0.000000
75	5%	2023.000000	1.800000e+05	175000.000000	100.000000
m	nax	2023.000000	3.040000e+07	450000.000000	100.000000

### Part I: Inferential Analysis 💆 💆

## Data Cleaning 🗸 🧽



After looking at some preliminary statistics for salary, we decided to undertake the following steps to clean our data and prepare it for use:

- 1. remote\_ratio column was divided by 100 to ensure scalability and correctness if the column were to be used as a numeric feature in a model.
- 2. work\_year was converted to string type as it is a categorical feature here and the interpretation of the 'mean' of the work\_year in the descriptive statistics is meaningless.

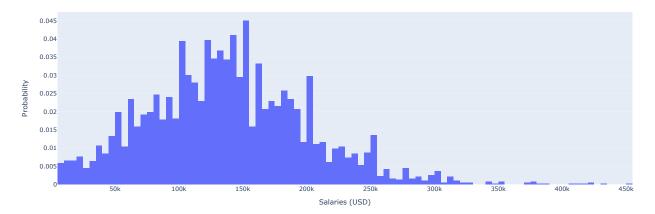
```
In [4]: salary['remote_ratio'].unique()
 Out[4]: array([100, 0, 50])
 In [5]: salary['remote_ratio'] = salary['remote_ratio'] / 100
 In [6]: salary.head()
             work_year experience_level employment_type
                                                                     job_title
                                                                              salary salary_currency salary_in_usd employee_residence remote_ratio company_location company_size
                                                                                                                                                                  FS
         0
                  2023
                                    SF
                                                      FT Principal Data Scientist
                                                                               80000
                                                                                                 FUR
                                                                                                             85847
                                                                                                                                   FS
                                                                                                                                                 1.0
          1
                 2023
                                    MI
                                                     СТ
                                                                  ML Engineer
                                                                               30000
                                                                                                LISD
                                                                                                             30000
                                                                                                                                   LIS
                                                                                                                                                 1.0
                                                                                                                                                                  US
                                                                                                                                                                                 S
          2
                  2023
                                    МІ
                                                      СТ
                                                                               25500
                                                                                                USD
                                                                                                             25500
                                                                                                                                   US
                                                                                                                                                 1.0
                                                                                                                                                                  US
                                                                                                                                                                                 S
                                                                  ML Engineer
         3
                  2023
                                    SE
                                                      FT
                                                                 Data Scientist 175000
                                                                                                 USD
                                                                                                            175000
                                                                                                                                   CA
                                                                                                                                                 1.0
                                                                                                                                                                  СА
                                                                                                                                                                                 М
                                    SE
                                                                 Data Scientist 120000
                                                                                                            120000
                                                                                                                                                 1.0
                                                                                                                                                                  СА
                                                                                                                                                                                 М
 In [7]: salary.shape
 Out[7]: (3755, 11)
 In [8]: salary['work_year'].dtype
Out[8]: dtype('int64')
 In [9]: salary['work_year'] = salary['work_year'].astype(str)
In [10]: salary.head()
          # print(salary.head().to_markdown(index=False))
             work_year experience_level employment_type
Out[10]:
                                                                     job_title
                                                                              salary_currency salary_in_usd employee_residence remote_ratio company_location company_size
         0
                 2023
                                    SE
                                                      FT Principal Data Scientist
                                                                              80000
                                                                                                 EUR
                                                                                                            85847
                                                                                                                                   ES
                                                                                                                                                 1.0
                                                                                                                                                                  ES
                                    МІ
                                                     СТ
                                                                                                                                                                  US
          1
                 2023
                                                                  ML Engineer
                                                                              30000
                                                                                                USD
                                                                                                            30000
                                                                                                                                   US
                                                                                                                                                 1.0
          2
                  2023
                                    MI
                                                      СТ
                                                                  ML Engineer
                                                                              25500
                                                                                                USD
                                                                                                             25500
                                                                                                                                   US
                                                                                                                                                 1.0
                                                                                                                                                                  US
                                                                                                                                                                                 S
         3
                                    SE
                                                      FT
                                                                                                                                   CA
                                                                                                                                                                                 М
                 2023
                                                                 Data Scientist 175000
                                                                                                USD
                                                                                                            175000
                                                                                                                                                 1.0
                                                                                                                                                                  CA
          4
                  2023
                                    SF
                                                      FT
                                                                 Data Scientist 120000
                                                                                                USD
                                                                                                            120000
                                                                                                                                   CA
                                                                                                                                                 1.0
                                                                                                                                                                  CA
                                                                                                                                                                                 М
```

#### EDA III


As part of our EDA we explored the relationship and the distribution of the following variables:

- 1. The Distribution of salary\_in\_usd
  - Observations: The distribution of salaries (USD) is fairly normal, and slightly right-skewed. It's a unimodal histogram with its peak at \$150,000-\$155,000
- 2. The Distribution of  $\verb|salary_in_usd|$  with respect to  $\verb|company_size|$ 
  - Observations: We can derive from the overlaid histogram that the medium companies are generally taller than the smaller and the larger companies. This is because the dataset contains more data-points for medium-sized companies. The boxplots accompanying the histograms confirm that the medium sized companies have more outliers and a higher median than the other two groups.
- 3. The average salary and salary\_in\_usd as work\_year increases
  - Observations: There is a steady increase in the mean salary\_in\_usd as the years increase. This could be explained by the usual rise in the general price level or inflation. However, the growth of the mean salary is quite peculiar. The salary column has a very high standard deviation, because of the disparities in the foreign exchange values of the different currencies. Thus, salary\_in\_usd provides a more standardized distribution of salaries.

### Distribution of Salaries (USD)




### Distribution of Salaries (USD)




```
In [14]: # fig.write_html('edaSecond.html', include_plotlyjs='cdn')
In [15]: fig = px.line(salary.groupby('work_year')['salary_in_usd', 'salary'].mean(), y=['salary_in_usd', 'salary'])
    fig.update_xaxes(title_text='Year')
    fig.update_yaxes(title_text='Salary (USD)')
    fig.show()

/var/folders/6g/cn7vfpqj1hq_5nknjxtfqhbh0000gn/T/ipykernel_61871/765794337.py:1: FutureWarning:
    Indexing with multiple keys (implicitly converted to a tuple of keys) will be deprecated, use a list instead.
```



```
In [16]: # fig.write_html('edaThird.html', include_plotlyjs='cdn')
In [17]: fig = px.scatter(salary[salary['salary_currency'] != 'USD'], y='salary', x='salary_in_usd', color='experience_level', hover_name='salary_currency',hover_data=['employee' fig.show()
```



## Hypothesis Testing 🧖

Null Hypothesis: The distribution of salaries of different company sizes is drawn from the same population, and any differences in the distribution are purely conincidental.

 $\textbf{Alternative Hypothesis:} \ \mathsf{The \ differences \ in \ salaries \ are \ not \ purely \ coincidental}$ 


Test Statistic: Mean s

p-value: 0.0

Conclusion: It is highly unlikely that the difference in means are purely coincidental. Thus, we reject the null hypothesis

```
x=1.25 * observed_tvd, showarrow=False, y=0.16)
fig.update_layout(xaxis_title='Difference in Means', yaxis_title='Probability', title_x=0.5)
```

#### Empirical Distribution of the Differences in Means



```
In [21]: # fig.write_html('hypothesis.html', include_plotlyjs='cdn')
In [22]: pval = np.mean(np.array(tvd) <= observed_tvd)
pval
Out[22]: 0.0</pre>
```

## Part II: Predictive Analysis 💇

Linear Regression -/ -/

Framing the Problem 🔣 😕

Prediction Problem: Predict the salaries of data scientists in USD

Response Variable: salary\_in\_usd

 $\textbf{Regressors} : \ \texttt{work\_year} \ , \ \texttt{employee\_type}$ 


Evaluation Metric: R2

```
In [23]: def plot_predictions(y_true, y_pred):
    fig = px.scatter(x=y_true, y=y_pred, labels={'x': 'Actual Values', 'y': 'Predicted Values'},
    title='Actual vs Predicted Values')
    fig.add_shape(type='line', x0=min(y_true), y0=min(y_true), x1=max(y_true), y1=max(y_true),
    line=dict(color='red', dash='dash'))
    fig.update_layout(title_x=0.5)
    return fig
```

## Baseline Model 🦑

R<sup>2</sup> score: 0.01


### Actual vs Predicted Values



```
In [25]: # fig.write_html('baseline.html', include_plotlyjs='cdn')
```

### Final Model 🏆

```
In [26]: fig = px.scatter(salary, x='work_year', y='salary_in_usd') fig.show()
```



R<sup>2</sup> score: 0.20


### Actual vs Predicted Values




### Actual vs Predicted Values


```
350k
      300k
Predicted Values
     200k
      150k
      100k
       50k
                                         50k
                                                                 100k
                                                                                        150k
                                                                                                               200k
                                                                                                                                       250k
                                                                                                                                                               300k
                                                                                                                                                                                      350k
                                                                                                                                                                                                              400k
                                                                                                       Actual Values
```

R<sup>2</sup> score: 0.20

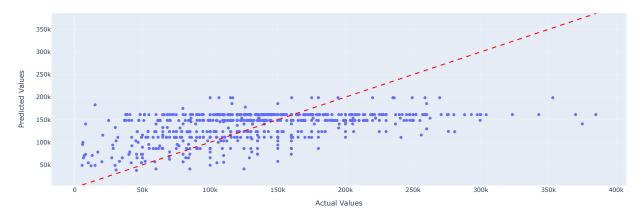


## Actual vs Predicted Values





## Actual vs Predicted Values




```
In [36]: # fig.write_html('final_mdl.html', include_plotlyjs='cdn')

Multiclass Classification
```

Prediction Problem: Predict the experience level of data scientists

Framing the Problem 🔣 🤥

```
Response Variable: experience_level
                       Features: employee type, company size
                       Evaluation Metric: Accuracy
                       Random Forest Classifier TT
In [37]: import pandas as pd
from sklearn.model_selection import train_test_split
                        from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score
                      # Prepare the data
X = salary[['employment_type', 'company_size']]
y = salary['experience_level']
                       # Perform one-hot encoding on the categorical features
                       X_encoded = pd.get_dummies(X)
                      # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)
                       # Create and train the random forest classifier
                       model = RandomForestClassifier()
                        model.fit(X_train, y_train)
                      # Make predictions on the test set
y_pred = model.predict(X_test)
                      # Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy: {:.2f}".format(accuracy))
precision = precision_score(y_test, y_pred, average=None)
                      # Display precision for each class
for class_label, class_precision in zip(model.classes_, precision):
    print("Precision for class {}: {:.2f}".format(class_label, class_precision))
                       precision
                  Accuracy: 0.67
Precision for class EN: 0.40
Precision for class EX: 0.00
Precision for class MI: 0.24
Precision for class SE: 0.69
                   /Users/shreya\_sudan/anaconda3/envs/projects 2023/lib/python 3.8/site-packages/sklearn/metrics/\_classification.py: 1469: \ Undefined Metric Warning: 1.0. \ An armonic of the package of 
                   Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
Out[37]: array([0.4
                                                           , 0.
                                                                                          , 0.24242424, 0.68863955])
In [38]: from sklearn.metrics import confusion_matrix
                      import plotly.graph_objects as go
import plotly.figure_factory as ff
                      # Assuming you have already trained your classification model and made predictions
y_true = y_test # True labels from the test set
y_pred = model.predict(X_test) # Predicted labels
                       cm = confusion_matrix(y_true, y_pred)
                        # Define the labels for the confusion matrix
                       labels = list(set(y_true))
                        # Create the confusion matrix figure using Plotly
                       fig = ff.create_annotated_heatmap(
                               z=cm,
x=labels,
                                  y=labels,
                                  colorscale='oryel',
                                  showscale=True,
                                  xgap=5.
                                 ygap=5
                       # Update the layout
                       fig.update_layout(
    title="Confusion Matrix",
    xaxis_title="Predicted Label",
    yaxis_title="True Label"
                        # Display the confusion matrix
                       fig.show()
```

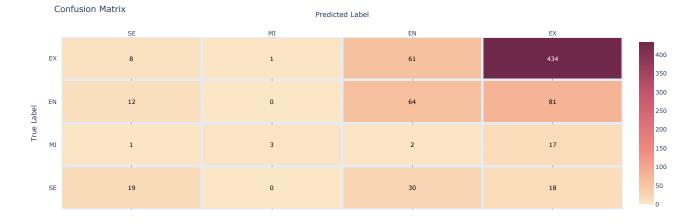


```
In [39]: # fig.write_html('randomforest.html', include_plotlyjs='cdn')
```

```
Decision Trees 🌳 🌳
```

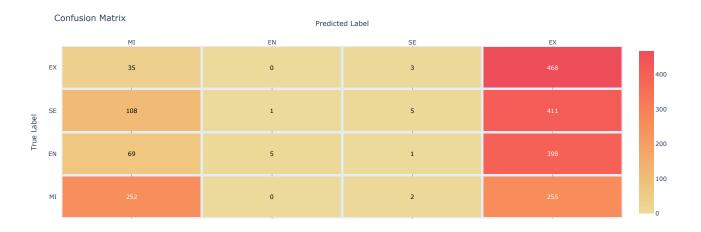
```
In [41]: import pandas as pd
                     from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
                     from sklearn.metrics import precision_score
                   # Prepare the data
X = salary[['employment_type', 'company_size']]
y = salary['experience_level']
                     # Perform one-hot encoding on the categorical features
                    X_encoded = pd.get_dummies(X)
                     # Split the data into training and testing sets
                    X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)
                   # Define the parameter grid for hyperparameter search
param_grid = {'max_depth': [3, 5, 7, 9]}
                    # Create the decision tree classifier
tree = DecisionTreeClassifier()
                     # Perform grid search with cross-validation
                   grid_search = GridSearchCV(tree, param_grid, cv=5)
grid_search.fit(X_train, y_train)
                    # Retrieve the best model
best_tree = grid_search.best_estimator_
                     # Make predictions on the test set using the best model
                    y_pred = best_tree.predict(X_test)
                     # Calculate precision for each class
                    precision = precision_score(y_test, y_pred, average=None)
                    # Display precision for each class
for class_label, class_precision in zip(best_tree.classes_, precision):
    print("Precision for class {}: {:.2f}".format(class_label, class_precision))
                     # Print the best hyperparameters
                   print("Best hyperparameters:", grid_search.best_params_)
accuracy = accuracy_score(y_test, y_pred)
                    print("Accuracy: {:.2f}".format(accuracy))
                 Precision for class EN: 0.40
                 Precision for class EX: 0.00
Precision for class MI: 0.24
Precision for class SE: 0.69
                 Best hyperparameters: {'max_depth': 5}
                 Accuracy: 0.67
                 /Users/shreya\_sudan/anaconda3/envs/projects 2023/lib/python 3.8/site-packages/sklearn/metrics/\_classification.py: 1469: \ Undefined Metric Warning: \ Analysis of the packages of the packag
                 Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
In [46]: import pandas as pd
                    from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
                     from sklearn.metrics import precision_score, recall_score
                   # Prepare the data
X = salary[['employment_type', 'company_size', 'salary_in_usd', 'job_title']]
y = salary['experience_level']
                     # Perform one-hot encoding on the categorical features
                    X_encoded = pd.get_dummies(X)
                    # Split the data into training and testing sets
                    X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)
                   # Define the parameter grid for hyperparameter search
param_grid = {'max_depth': [3, 5, 7, 9]}
                    # Create the decision tree classifier
                    tree = DecisionTreeClassifier()
```

```
# Perform grid search with cross-validation
grid_search = GridSearchCV(tree, param_grid, cv=5)
                grid_search.fit(X_train, y_train)
                # Retrieve the best model
                best_tree = grid_search.best_estimator_
               # Make predictions on the test set using the best model
y_pred = best_tree.predict(X_test)
               # Print the best hyperparameters
print("Best hyperparameters:", grid_search.best_params_)
                accuracy = accuracy_score(y_test, y_pred)
                print("Accuracy: {:.2f}".format(accuracy))
             Best hyperparameters: {'max_depth': 9}
In [47]: precision = precision_score(y_test, y_pred, average=None)
                # Display precision for each class
                # Display precision for each class
for class_label, class_precision in zip(best_tree.classes_, precision):
    print("Precision for class {}: {:.2f}".format(class_label, class_precision))
recall = recall_score(y_test, y_pred, average=None)
# Display precision for each class
for class_label, class_precision in zip(best_tree.classes_, recall):
    print("Recall for class {}: {:.2f}".format(class_label, class_precision))
             Precision for class EN: 0.47
Precision for class EX: 0.75
Precision for class MI: 0.41
Precision for class SE: 0.79
            Recall for class EN: 0.28
Recall for class EX: 0.13
Recall for class MI: 0.41
Recall for class SE: 0.86
In [48]: import plotly.figure_factory as ff
from sklearn.metrics import confusion_matrix
                # Assuming you have already trained your classification model and made predictions
y_true = y_test # True labels from the test set
y_pred = best_tree.predict(X_test) # Predicted labels
                # Create a confusion matrix
                cm = confusion_matrix(y_true, y_pred)
                # Define the labels for the confusion matrix
labels = list(set(y_true))
                   Create the confusion matrix figure using Plotly
                fig = ff.create_annotated_heatmap(
   z=cm,
   x=labels,
                       v=labels.
                      colorscale='burgyl',
showscale=True,
                       xgap=5,
                       ygap=5
                # Undate the lavout
               fig.update_layout(
    title="Confusion Matrix",
    xaxis_title="Predicted Label",
    yaxis_title="True Label"
                # Display the confusion matrix
                #burgyl, deep, magenta, matter, mint, oryel, peach, pinkyl, purp, purpor, redor, sunset?, sunsetdark, teal and tealgrn???,
```




```
In [49]: # fig.write_html('decisiontree.html', include_plotlyjs='cdn')
# ['aggrnyl', 'agsunset', 'algae', 'amp', 'armyrose', 'balance',
# 'blackbody', 'bluered', 'blues', 'blugrn', 'bluyl', 'brbg',
# 'brwnyl', 'bupn', 'bupn', 'burg', 'burgyl', 'civ'dis', 'curl',
# 'darkmint', 'deep', 'delta', 'dense', 'earth', 'edge', 'electric',
# 'emrld', 'fall', 'geyser', 'gnbu', 'gray', 'greens', 'greys',
# 'haline', 'hot', 'hsv', 'ice', 'icefire', 'inferno', 'jet',
```

```
'magenta', 'magma', 'matter', 'mint', 'mrybm', 'mygbm', 'oranges',
'orrd', 'oryel', 'oxy', 'peach', 'phase', 'picnic', 'pinkyl',
'piyg', 'plasma', 'plotly3', 'portland', 'prgn', 'pubu', 'pubugn',
'puor', 'purd', 'purp', 'purples', 'purpor', 'rainbow', 'rdbu',
'rdgy', 'rdpu', 'rdylbu', 'rdylgn', 'redor', 'reds', 'solar',
'spectral', 'speed', 'sunset', 'sunsetdark', 'teal', 'tealrn',
'tealrose', 'tempo', 'temps', 'thermal', 'tropic', 'turbid',
'turbo', 'twilight', 'viridis', 'ylgn', 'ylgnbu', 'ylorbr',
'vlorrd',
                                               'ylorrd']
In [50]: from imblearn.over_sampling import SMOTE
from sklearn.preprocessing import StandardScaler
                 import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score
                 # Prepare the data
X = salary[['employment_type', 'company_size']]
y = salary['experience_level']
                  # Perform one-hot encoding on the categorical features
                 X_encoded = pd.get_dummies(X)
                 smote = SMOTE()
                 X_resampled, y_resampled = smote.fit_resample(X_encoded, y)
                 # Scale the features using StandardScaler
scaler = StandardScaler()
                 X_scaled = scaler.fit_transform(X_resampled)
                 # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_resampled, test_size=0.2, random_state=42)
                 # Create and train the random forest classifier
                                 RandomForestClassifier()
                 model.fit(X_train, y_train)
                 # Make predictions on the test set
                 y_pred = model.predict(X_test)
                # Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy: {:.2f}".format(accuracy))
               Accuracy: 0.36
In [51]: precision = precision_score(y_test, y_pred, average=None)
                 # Display precision for each class
for class_label, class_precision in zip(best_tree.classes_, precision):
    print("Precision for class {}: {:.2f}".format(class_label, class_precision))
recall = recall_score(y_test, y_pred, average=None)
                 # Display precision for each class
for class_label, class_precision in zip(model.classes_, recall):
    print("Recall for class {}: {:.2f}".format(class_label, class_precision))
              Precision for class EN: 0.54
Precision for class EX: 0.83
Precision for class MI: 0.45
Precision for class SE: 0.31
              Recall for class EN: 0.50
Recall for class EX: 0.01
Recall for class MI: 0.01
Recall for class SE: 0.92
In [52]: from sklearn.metrics import confusion matrix
                 import plotly.graph_objects as go
import plotly.figure_factory as ff
                 # Assuming you have already trained your classification model and made predictions
y_true = y_test # True labels from the test set
y_pred = model.predict(X_test) # Predicted labels
                  # Create a confusion matrix
                 cm = confusion_matrix(y_true, y_pred)
                  # Define the labels for the confusion matrix
labels = list(set(y_true))
                  # Create the confusion matrix figure using Plotly
                  fig = ff.create_annotated_heatmap(
                         z=cm.
                         v=labels.
                         colorscale='oryel',
showscale=True,
                         xgap=5,
ygap=5
                  # Update the layout
                 # Opdate the Layout
fig.update_layout(
    title="Confusion Matrix",
    xaxis_title="Predicted Label",
    yaxis_title="True Label"
                  # Display the confusion matrix
                 fig.show()
```



```
In [56]:
import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import precision_score
             # Prepare the data
X = salary[['employment_type', 'company_size', 'salary_in_usd', 'job_title']]
y = salary['experience_level']
             # Perform one-hot encoding on the categorical features
             X_encoded = pd.get_dummies(X)
              smote = SMOTE()
             X_resampled, y_resampled = smote.fit_resample(X_encoded, y)
             # Scale the features using StandardScaler
scaler = StandardScaler()
              X_scaled = scaler.fit_transform(X_resampled)
             # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_resampled, test_size=0.2, random_state=42)
             # Define the parameter grid for hyperparameter search
param_grid = {'max_depth': [3, 5, 7, 9]}
             # Create the decision tree classifier
             tree = DecisionTreeClassifier()
              # Perform grid search with cross-validation
             grid_search = GridSearchCV(tree, param_grid, cv=5)
grid_search.fit(X_train, y_train)
             # Retrieve the best model
best_tree = grid_search.best_estimator_
             # Make predictions on the test set using the best model
             y_pred = best_tree.predict(X_test)
             # Print the best hyperparameters
print("Best hyperparameters:", grid_search.best_params_)
             accuracy = accuracy_score(y_test, y_pred)
print("Accuracy: {:.2f}".format(accuracy))
            Best hyperparameters: {'max_depth': 9}
           Accuracy: 0.67
In [57]: import plotly.figure_factory as ff
from sklearn.metrics import confusion_matrix
             # Assuming you have already trained your classification model and made predictions
y_true = y_test # True labels from the test set
y_pred = best_tree.predict(X_test) # Predicted labels
             # Create a confusion matrix
             cm = confusion_matrix(y_true, y_pred)
             # Define the labels for the confusion matrix
labels = list(set(y_true))
              # Create the confusion matrix figure using Plotly
              fig = ff.create_annotated_heatmap(
                   z=cm.
                   x=labels,
y=labels,
                   colorscale='burgyl',
showscale=True,
                   xgap=5,
ygap=5
             fig.update_layout(
    title="Confusion Matrix",
    xaxis_title="Predicted Label",
    yaxis_title="True Label"
             # Display the confusion matrix
             fig.show()
```

#burgyl, deep, magenta, matter, mint, oryel, peach, pinkyl, purp, purpor, redor, sunset?, sunsetdark, teal and tealgrn???,
#



```
In [58]: precision = precision_score(y_test, y_pred, average=None)

# Display precision for each class
for class_label, class_precision in zip(best_tree.classes_, precision):
    print("Precision for class {}: {:.2f}".format(class_label, class_precision))
    recall = recall_score(y_test, y_pred, average=None)

# Display precision for each class
for class_label, class_precision in zip(best_tree.classes_, recall):
    print("Recall for class {}: {:.2f}".format(class_label, class_precision))

Precision for class EN: 0.71
Precision for class EN: 0.65
Precision for class SE: 0.65
Precision for class SE: 0.69
Recall for class EN: 0.85
Recall for class EN: 0.85
Recall for class EN: 0.85
Recall for class SE: 0.62

In [59]: # fig.write_html('final_classification.html', include_plotlyjs='cdn')
```

# Conclusion 🚟 🕮